

Notation

This section discusses notation used in this specification.

All operations in this document are implicitly performed modulo 2^{32} . We use standard mathematical notation for addition, subtraction, multiplication, and exponentiation. Division always denotes integer division, i.e. any remainder is dropped.

We also use the following operators:

- $x \ll y$ denotes shifting x to the left y bits, i.e. $x \ll y = x2^y$
- $x \gg y$ denotes a *logical* right shift – it shifts x to the right by y bits, i.e. $x \gg y = \frac{x}{2^y}$
- $\text{ROTL}^n(x)$ rotates x n bits to the left, i.e. $\text{ROTL}^n(x) = (x \ll n) + (x \gg (32 - n))$

The RRS Rolling Checksums

The **rrs** family of checksums were first used in **rsync**, and later in **bup** and **perkeep**. **rrs** was originally inspired by the **adler-32** checksum. The name **rrs** was chosen for this specification, and stands for **rsync rolling sum**.

Definition

A concrete **rrs** checksum is defined by the parameters:

- M , the modulus
- C , the character offset
- R , the rotation.

Given a sequence of bytes X_0, X_1, \dots, X_N and a choice of M and C , the **rrs** hash of the sub-sequence X_k, \dots, X_l is $s(k, l)$, where:

$$a(k, l) = (\sum_{i=k}^l (X_i + C)) \bmod M$$

$$b(k, l) = (\sum_{i=k}^l (l - i + 1)(X_i + C)) \bmod M$$

$$s(k, l) = \text{ROTL}^R(a(k, l) + 2^{16}b(k, l))$$

RRS0

The concrete hash called **rrs0** uses the values:

- $M = 2^{16}$
- $C = 31$
- $R = 0$

`rrs0` is used by current versions of `librsync` as of August 2020. Note that the hash in the `rsync` documentation is not `rrs0`; that hash uses $C = 0$.

RRS1

The concrete hash called `rrs1` uses the values:

- $M = 2^{16}$
- $C = 31$
- $R = 16$

`rrs1` is used by both Bup and Perkeep, and implemented by the go package `go4.org/rollsum`.

Implementation

Rolling

`rrs` is a family of *rolling* hashes. We can compute hashes in a rolling fashion by taking advantage of the fact that:

$$a(k+1, l+1) = (a(k, l) - (X_k + C) + (X_{l+1} + C)) \bmod M$$

$$b(k+1, l+1) = (b(k, l) - (l - k + 1)(X_k + C) + a(k+1, l+1)) \bmod M$$

So, a typical implementation will work like:

- Keep X_k, \dots, X_l in a ring buffer.
- Also store $a(k, l)$ and $b(k, l)$.
- When X_{l+1} is added to the hash:
 - Dequeue X_k from the ring buffer, and enqueue X_{l+1} .
 - Use X_k, X_{l+1} , and the stored $a(k, l)$ and $b(k, l)$ to compute $a(k+1, l+1)$ and $b(k+1, l+1)$. Then use those values to compute $s(k+1, l+1)$ and also store them for future use.

Choice of M

Choosing $M = 2^{16}$ has the advantages of simplicity and efficiency, as it allows $s(k, l)$ to be computed using only shifts and bitwise operators; in C:

```
uint32_t s = a | (b << 16);
```